arXiv:2611.14117v1 [q-fin.TR] 10 Nov 2026

Fine-Tuning an Order-Flow Foundation Model for Sub-Second Forecasting

Fedor Shabashev

November 11, 2026

Abstract

Generative models have become the standard approach in language and vision. These models are typically pretrained first and then fine-tuned. Recent work demonstrates how to pretrain an autoregressive generative model on Level-3 order-flow data. However, to our knowledge, no one has fine-tuned such a model yet for downstream tasks.

Here we first pretrain a 300M-parameter generative model on Level-3 order-flow data collected from two crypto exchanges, Bitstamp and Bitfinex. Then we fine-tune the pretrained model to predict the microprice return and its realized volatility. We show that our fine-tuned model outperforms classical baselines, such as gradient boosting on handcrafted features, in predicting microprice returns over sub-second horizons. Furthermore, the predictive ability of the model transfers zero-shot to assets it was never fine-tuned on. Due to the short-term nature of these forecasts, the predicted microprice moves are smaller than the bid-ask spread plus fees, so the natural applications are in execution and market making.

These results indicate that a pretrained generative model can be turned into a competitive short-horizon price forecaster via fine-tuning.

1 Introduction

Level-3 market data is a sequence of orders that market participants send to the exchange. Recent work [Kawawa-Beaudan et al., 2026, Sood et al., 2026] shows that a generative model can be trained to autoregressively predict the next event in the order sequence, one event at a time. For most financial markets, including the crypto markets we model here, the orders arrive at high frequency, often hundreds of orders per second. At the same time submitting your own order to an exchange might take more than 100 ms and receiving data from the exchange takes tens of milliseconds as well [Jha et al., 2020, Albers et al., 2025, Foley et al., 2023]. Hence a market participant will want to know the market prices hundreds of milliseconds in advance, which is a much longer horizon than just one order ahead. So a generative model predicting the next order by itself cannot be used directly to answer practical questions such as "is it a good idea to submit a limit order at bid right now", as by the time the order arrives to the exchange it will be outside of the modeling horizon of the generative model.

To convert the order flow into a form a neural network can consume, we tokenize the orders [Kawawa-Beaudan et al., 2026, Sood et al., 2026]. In our approach we normally convert one order into one token, so ten buy limit orders will be converted into ten tokens in the sequence. Crypto exchanges also allow their participants to modify resting orders by changing the price or size. These update orders we convert into two separate orders, a delete and an add, with the same timestamp.

On the tokenized order flow, 2.8 billion training tokens, we train an autoregressive decoder-only Transformer model with 300M parameters. The model is trained to predict the next token with a

cross-entropy loss.

We then take the generative model and fine-tune it via LoRA to predict the first and second moments of the future microprice on the BTC/USD pair. The microprice prediction is done for several horizons: 200 ms, 500 ms, 1, 2, 5, 10 and 30 seconds. For each of these horizons our model predicts the change in microprice and its volatility. We then benchmark the fine-tuned model against classical benchmarks such as gradient boosting and GARCH. We demonstrate that our fine-tuned model outperforms all classical benchmarks in predicting the microprice change direction at the 200 ms and 500 ms horizons, and outperforms GARCH on volatility in four of five evaluation windows. Furthermore, we demonstrate that the fine-tuned model can be applied zero-shot to new assets it has not been fine-tuned on and is capable of beating basic baselines there.

Contributions.

1. To our knowledge, the first fine-tuning of an autoregressive order-flow foundation model for predicting future price and volatility. We fine-tuned such a model and benchmarked it against classical baselines: order-book imbalance, order-flow imbalance, gradient boosting on those features, GARCH and persistence. We found that the fine-tuned model outperforms the baselines for predicting future price at sub-second horizons (Tables 4–7).

2. Zero-shot asset transfer. A generative model fine-tuned to predict the future microprice for BTC (Bitcoin) outperforms basic baselines when it comes to predicting the future microprice for SOL (Solana) and ETH (Ethereum).

3. We found that the fine-tuned model performs well at sub-second horizons, but for horizons longer than a few seconds the classical methods match or beat it, which highlights the limitations of the current approach.

2 Related Work

Supervised LOB forecasting. A long line of work predicts short-horizon price moves directly from the book with task-specific supervised models. Cont et al. [2014] show that order-flow imbalance (OFI)—net signed size at the best quotes—linearly predicts contemporaneous price change, and Kolm et al. [2023] extend multi-level OFI into a deep, multi-horizon alpha signal; Zhang et al. [2019] (DeepLOB) train convolutional-recurrent networks on raw book snapshots for direction, and Sirignano and Cont [2019] show that deep networks on LOB states learn price-formation features that are largely universal across assets. Our directional baselines—multi-level order-book imbalance and signed trade-flow, plus multi-level OFI at matched lookbacks as the sequence-aware control—are the linear members of this family. The distinction from our work is that these models are trained from scratch on labeled targets, whereas we fine-tune a representation pretrained generatively on unlabeled order flow.

Generative LOB models. A separate thread models the *distribution* of order flow rather than a supervised target. Nagy et al. [2023] introduce a token-level autoregressive generative model of message flow; Kawawa-Beaudan et al. [2026], Sood et al. [2026] (TradeFM) scale this into a foundation model with the one-token-per-trade tokenization we adopt, and Nagy et al. [2025] (LOB-Bench) provide the K–S / 1-Wasserstein marginal protocol we use to assess generative fidelity. Li et al. [2025] (MarS) build an order-level generative foundation model of comparable ambition and do obtain forecasts from it—but *through simulation*: they roll the generator forward 128 times from

each starting point and aggregate the trajectories into a minute-scale up/down/flat class, leaving the backbone's weights untouched and contrasting their approach with "task-specific forecasting" precisely because no task-specific training occurs. In our model that route would pay the price Section 5 measures: an open-loop rollout accumulates the deep-sweep volatility error. Reading the forecast off the representation instead, as we do, avoids it. These models optimize distributional realism and, in the case of TradeFM, explicitly leave task-specific fine-tuning as future work—the gap this paper addresses. Linna et al. [2025] (LOBERT) trained a 1.1M-parameter BERT-like encoder model on four NASDAQ stocks and fine-tuned it for mid-price direction; the autoregressive backbones above have not been fine-tuned. Acero [2026] pretrains GPT-style models of up to 32M parameters on synthetic order flow from simulators with known parameters and continues next-token training on one day of NASDAQ messages, where the synthetic initialization beats from-scratch training on held-out tickers; the objective remains next-token loss, so no forecasting task is fine-tuned. The classical zero-intelligence generator [Smith et al., 2003], which we use as a baseline, Hawkes-process models of order flow [Bacry et al., 2015], and the canonical return stylized facts [Cont, 2001] set what a faithful generator must reproduce.

Time-series foundation models. Chronos [Ansari et al., 2024], TimesFM [Das et al., 2024], and Moirai [Woo et al., 2024] pretrain on broad corpora of sampled time series and transfer zero-shot or with light fine-tuning. They operate on regularly-sampled scalar or vector series, not event-level Level-3 order flow, so adapting them to the LOB—where the signal lives in the irregular *sequence* of individual book events—is a different problem from fine-tuning a native LOB generative model, which is what we study. This is the sense in which the present fine-tuning is new: supervised LOB predictors are trained from scratch, generative LOB models target distributional fidelity, and time-series foundation models are not native to order flow.

Microstructure baselines and targets. Our volatility baselines are realized-volatility persistence and GARCH [Bollerslev, 1986] (HAR [Corsi, 2009] is the standard long-memory alternative, not run here); our forecast target is the microprice, the high-frequency fair-value estimator of Stoikov [2018]. Together with OFI/imbalance for direction, these are the baselines the fine-tuned model is measured against.

3 Data and Tokenization

Data. We collect every Level-3 (L3) order-book event (adds, changes, cancels, and trades) from two exchanges (Bitfinex, Bitstamp). The main model's corpus spans roughly two weeks (20 September to 2 October 2026); the recency control of Section 4 uses later, longer collections from the same feeds.

Event rates. Crypto L3 is higher-frequency per symbol than the equity ITCH feeds used by prior LOB models, and steadier (Fig. 1). On Bitstamp, BTC/USD runs a median of 199 events/s; measured from NASDAQ TotalView-ITCH 5.0 with the same event definitions, the busiest equities run 3–4× slower in the median (SPY 62/s, AAPL 52/s), though spikier—single-second microbursts (AAPL >28,000/s) that crypto never shows (BTC max ∼1,000/s). Two structural differences shape the tokenizer: crypto venues stamp on a millisecond grid, collapsing 32% of consecutive events onto an identical timestamp (vs 5–8% on ITCH's nanosecond grid), and trade continuously rather than in a 6.5 h session.

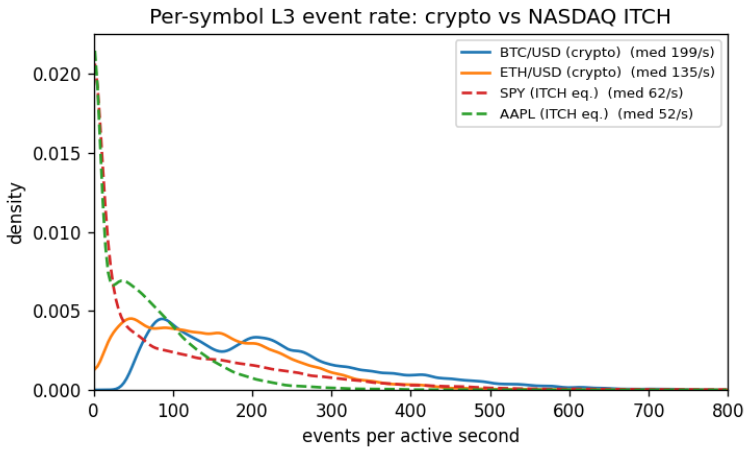

Figure 1: Per-symbol event-rate density, crypto L3 vs NASDAQ ITCH equities (both measured with our tokenizer's event definitions). Crypto (solid) sustains a markedly higher per-second rate than even the busiest equities (dashed).

Tokenizer. Each add or cancel becomes one token with fields (action, side, δp, $\log v$, Δt) at radices $(2, 2, 16, 16, 16)$, a 16,384-element vocabulary, exactly as TradeFM. Trades do not receive a token of their own: the stream is add/cancel only, the liquidity an execution consumes appears as the removal of the resting order, and trades enter the representation through the trade-based EW-VWAP mid against which δp is measured. Δt is the backward-looking inter-arrival time since the preceding event. Crypto required three departures from the equities recipe:

- **Reserved Δt=0 bin.** Order activity is bursty at the millisecond: 32% of consecutive events carry an identical timestamp—predominantly fleeting limit orders posted and canceled within one tick, plus orders batched in a single exchange message. We therefore reserve inter-arrival bin 0 for Δt=0 and quantile-bin only the strictly-positive tail into the remaining bins.
- **Reprice splitting.** A Bitfinex `change` (an order moving price) is emitted as a cancel of the old resting order followed by an add of the new one, so the token stream is pure add/cancel.
- **Pooled calibration, frozen edges.** Bin edges are fit once on a pooled sample of both venues and frozen, so a single token id denotes the same event across venues and across retrains.

Continuous fields are scale-invariant ratios: δp is signed basis points from the current mid, $\log v$ is $\log(1 + \text{size})$, and the mid is an exponentially-weighted VWAP of executed trades: a running average of trade prices, each weighted by its size, with a 30 s exponential half-life so recent and larger trades dominate. (We use a trade-VWAP mid rather than the book midpoint because it is robust to a one-sided or momentarily empty book.) Table 1 traces a run of consecutive BTC/USD events through the tokenizer.

4 Model and Training

We train a stock `LlamaForCausalLM` [Touvron et al., 2023] (RoPE, GQA, SwiGLU, RMSNorm), context 1024 tokens, our main model being 300M parameters (hidden 1152, 18 layers), for 4 epochs over 2.76B training tokens. The held-out split is a temporally-later chunk (412M tokens), scored under one fixed protocol (500 batches, fixed seed) for all cross-model comparisons.

Table 1: Eight consecutive Bitstamp BTC/USD book events and the tokens they produce (26 October 2026; EW-VWAP mid 79,932.45). δp is signed basis points from the mid, $\log v = \log(1 + \text{size})$, and Δt the gap since the preceding kept event. The bin tuple is (action, side, δp, $\log v$, Δt) with action 0 = add, 1 = cancel and side 0 = bid, 1 = ask, and the id is its mixed-radix value $8192a + 4096s + 256p + 16v + t$. The bid 10% below the mid falls in the lowest δp bin; the sixth row lands in the reserved Δt=0 bin. Bitstamp stamps on a millisecond grid, and the quantile edges at 0.9999 and 1.0002 ms bracket it, so a 1 ms gap lands in bin 4 or 5 depending on floating-point rounding of the timestamp difference; the edges are frozen, so this redundancy is left in place.

event	side	price	size	Δt (ms)	δp (bp)	$\log v$	(a, s, p, v, t)	id
add	bid	71,939.21	0.0014	3	−1000.0	0.001	(0,0,0,0,7)	7
cancel	bid	71,939.21	0.0014	13	−1000.0	0.001	(1,0,0,0,10)	8202
cancel	ask	79,940.91	0.0625	1	+1.06	0.061	(1,1,8,2,4)	14372
cancel	ask	79,932.46	0.2085	4	+0.00	0.189	(1,1,7,2,7)	14119
cancel	ask	79,932.46	0.1000	1	+0.00	0.095	(1,1,7,2,5)	14117
cancel	ask	79,932.46	0.1000	0	+0.00	0.095	(1,1,7,2,0)	14112
cancel	ask	79,950.00	0.7149	1	+2.20	0.539	(1,1,8,3,4)	14388
add	ask	79,958.00	0.1200	1	+3.20	0.113	(0,1,9,2,4)	6436

Scaling. The regime is data-constrained (Table 2): loss falls steadily with corpus size through 2.76B tokens, whereas enlarging the model 1.53× at fixed data (300M→460M) buys only −0.015 nats. We do not read a scaling law from the data column. The corpora grew by accretion, so each larger corpus is also a more recent one and the holdout is later than all of them, and a control that holds the corpus end date fixed while varying only its size (Fig. 2) finds that recency accounts for most of the slope: on a later holdout, a corpus ending 15 days closer to it is worth ∼0.5 nats at equal data budget, while 3.9× more data at a fixed end date is worth ∼0.16. Downstream the order reverses: the same two contrasts move the fine-tuned model's 0.2 s directional correlation by +0.005 (recency) and +0.012 (size), one backbone pair per contrast and three fine-tuning seeds per backbone (seed s.d. 0.003–0.007; the size contrast is seed-disjoint at all three sub-second horizons, the recency contrast only at 0.2 s). Data still dominates model size in this range, by less than the raw column suggests.

Table 2: Next-token loss (nats) on the held-out chunk, fixed protocol. Left: data scaling at 300M. Right: the model-constrained test at fixed 2.76B.

data scaling (300M)		
train tokens	NLL	
1.23B	2.742	
1.63B	2.647	
2.26B	2.452	
2.76B	**2.309**	
model scaling (2.76B)		
300M	2.309	
460M	2.294	(−0.015)

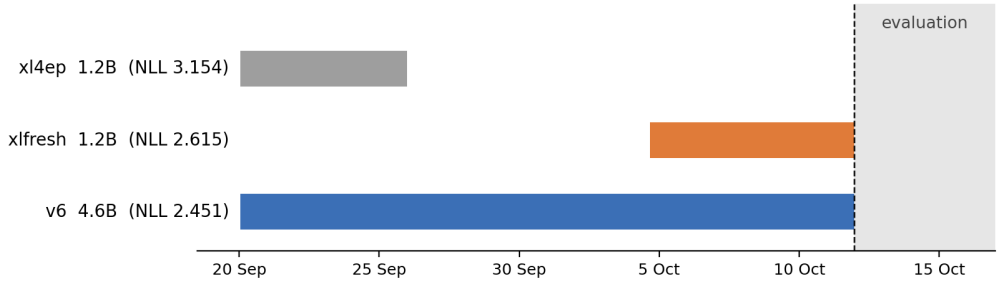

Figure 2: Controlling for corpus recency. Three pretraining corpora on a date axis. *xl4ep* and *xlfresh* have the same data budget and end 15 days apart; *xlfresh* and *v6* end at the same instant and differ 3.9× in size. Each pair therefore varies one factor. All three are scored in the same later period (shaded): next-token loss on a 12–13 October chunk after the cutoff, and 0.2 s direction after LoRA fine-tuning on 13–15 October with the 20k-step recipe of Section 6, three seeds per backbone, scored on the 33.5-hour window of 15–16 October, which is disjoint from the fine-tuning set and begins seven hours after it ends. Recency moves the loss far more than size does (−0.54 vs −0.16 nats), but size moves the downstream forecast more than recency does (+0.012 vs +0.005 correlation). One backbone pair per contrast. *v6* serves only as the control's large arm, not as a candidate main model: its corpus runs to 11 October and so overlaps the fine-tuning data and the 8 and 11 October evaluation windows of Section 6 (the 5 October holdout is excluded from it).

5 Generative fidelity

Before asking what the representation carries, we check that the pretrained model is a faithful generator. We label the checks by market-data level: L2 denotes book-level quantities (spread, depth, imbalance) and L1 the price path (returns), with L3 the order-flow events themselves. A rollout samples tokens autoregressively at temperature 1; each token is de-quantized to an event and applied to a matching engine, where a marketable add fills against the resting book and moves the EW-VWAP mid, a non-marketable add rests, and a cancel removes the nearest resting order (token ids carry no order identity, as in TradeFM). Feeding real tokens through the same engine reconstructs a sensible two-sided book (spread ∼1–2 bps). Against the real continuation of the same seed windows, and through the identical engine, we compare the model's rollouts and a *zero-intelligence* (ZI) stream that resamples each field independently from its real marginal [Smith et al., 2003], matching every marginal but destroying joint and temporal structure. Following LOB-Bench [Nagy et al., 2025] and TradeFM, distributions are compared by Kolmogorov–Smirnov distance (D_n) and 1-Wasserstein (W_1), with no pairing and no point forecast.

Book factors. Table 3 shows the model matches the real book's emergent factors as tightly as TradeFM's original equities model and beats ZI on three of the four (spread, depth imbalance, resting-order count—ZI's book collapses to 28 resting orders), ZI edging it only on touch imbalance. On the joint structure of the token fields the model is indistinguishable from a second real sample (Jensen–Shannon divergence at the two-sample floor on three field joints, where ZI sits 5–9× above it).

Returns and stylized facts. The price process that emerges from the model's flow reproduces the stylized facts—near-zero return autocorrelation and slowly decaying $|r|$-autocorrelation—which ZI lacks entirely, and its return distributions are closer to real than ZI's at every sub-second horizon

Table 3: L2 emergent factors, K-S distance to real (64 rollouts × 8192 events). TFM column is the original equities model for reference. ZI's thin book is also lopsided (median depth imbalance 0.50 against 0.00 real), which its K-S of 0.56 reflects.

factor	median real	median model	median ZI	K-S model	K-S ZI	TFM (equities)
spread (bps)	1.84	2.22	3.16	**0.20**	0.53	0.238
touch imbalance	0.111	0.121	0.000	0.16	0.09	0.142
depth imbalance (top-10)	0.004	−0.008	0.499	**0.19**	0.56	0.169
resting order count	219	219	28	**0.04**	0.79	—

(K-S 0.02–0.06 against 0.11–0.33; Fig. 3). The typical rollout is calibrated to within ∼1.2× of real volatility, but rare rollouts in which persistent same-side flow sweeps ∼4× deeper into the book than real flow does (median crossing depth 5.1 vs 1.4 bps) inflate the pooled std-based volatility ratio to ∼15×, and the inflation grows with rollout depth past about one context window (∼5 s). Two consequences matter for what follows. First, K-S misses this tail entirely: most sub-second returns are an exact-zero atom, so the maximum CDF gap sits at the mode, whereas W_1 registers it; we therefore report W_1 and realized volatility rather than K-S alone. Second, across five checkpoints spanning 0.5 nats of loss the deep-sweep frequency shows no monotone relation to loss: next-token loss and open-loop price realism are decoupled, so reading a forecast off the representation, as we do next, rather than off a rollout, avoids a cost that scale does not remove.

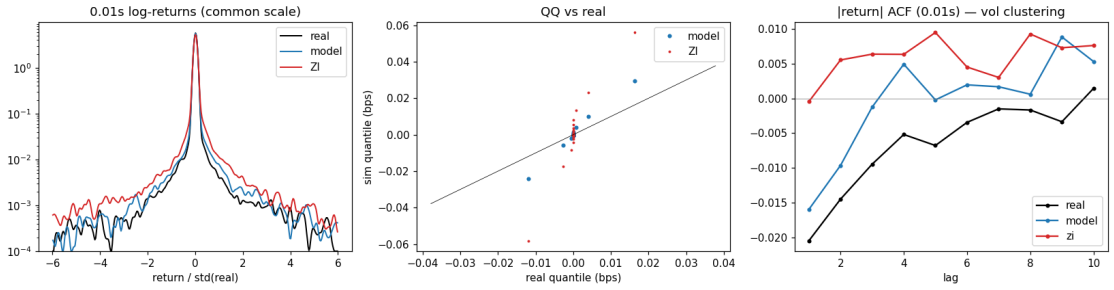

Figure 3: Sub-second L1: return density on a common scale (all variants divided by the real std), QQ vs real, and $|r|$-autocorrelation. Model (blue) tracks real (black) in shape and reproduces volatility clustering; ZI (red) is visibly fatter. The common scale prevents a single rollout's deep-sweep outlier from rescaling an entire distribution.

6 Fine-tuning for the first two moments

We fine-tune the pretrained model to forecast, from the token context, the first moment (signed mid return r_h; the mid is the microprice throughout, defined below) and second moment (realized volatility rv_h) of the near-future price across horizons from 200 ms to 30 s. The question is whether the pretrained representation earns its keep against standard microstructure baselines—so we benchmark, test for leakage, and measure cross-asset transfer.

Supervised dataset. We re-tokenize a slice of one asset (bitstamp:btcusd) with the frozen edges—so tokens match the pretrained corpus—and compute targets from the reconstructed *microprice* $(P_bQ_a+P_aQ_b)/(Q_a+Q_b)$ anchored at each token's *true exchange timestamp*. The book is

rebuilt from L3 and *re-seeded from a fresh snapshot every ∼minute* so reconstruction error cannot accumulate; the reconstructed microprice matches polled snapshots to 0.0 bps (median and p90 over sampled instants). The microprice (not the sticky price-midpoint) is used because the BTC touch price moves rarely, giving an otherwise-degenerate vol target. $r_h = \log m(t{+}h) - \log m(t)$ and $rv_h = \sqrt{\sum(\Delta \log m)^2}$ over $(t, t{+}h]$, with the base mid taken *after* the current event (the model has already seen it), so r_h excludes the token's own price impact—conservative, not leaky. An independent per-position recompute matches the vectorized targets to 10^{-6} bps and tokens re-tokenize deterministically.

Token-time alignment. Each token's forward target must be anchored at its true exchange time. Reconstructing token times by *cumulative-summing* the tokenizer's continuous Δt is not equivalent: dropped warm-up and invalid-price events advance the clock without emitting a token, so a cumsum clock runs ahead with accumulating drift—here ∼2.25 s, about 225% of a 1 s horizon—which shifts every "forward" window into the token's own past and overlaps its context, corrupting the target and inflating apparent skill. We therefore emit the true per-token exchange timestamp and validate the resulting targets by deterministic re-tokenization, an independent per-position recompute, and snapshot ground truth. For an ultra-short-horizon forward target, how a token is placed in time is load-bearing.

Head and transfer. The last hidden state feeds a small head $(\mu_h, \log \sigma_h)$ per horizon, trained with a robust decoupled loss $\text{Huber}(\mu, r) + \text{Huber}(\log \sigma, \log rv)$. We use Huber rather than a heteroscedastic Gaussian NLL because a naive NLL run collapsed the mean to a near-constant, the documented failure of Seitzer et al. [2022] in which the optimizer suppresses μ by inflating σ; the decoupled loss avoids it by construction. We compare a *frozen probe* (head only) against *LoRA* [Hu et al., 2022] (rank-16 attention adapters, 0.7% of weights); the tables label the pretrained backbone FM, for foundation model. LoRA beats the probe where there is signal to extract, most at the shortest horizons (+0.11 direction correlation at 200 ms versus +0.01 at 5 s): adaptation surfaces signal the frozen features do not, concentrated exactly where the model's edge lives. Tables 4 and 5 use 1,500-step constant-LR fine-tunes; the seeded analyses later in this section use the 10k-step cosine recipe, which reproduces the sub-second rows within the seed spread. LoRA: rank 16, α=32, dropout 0.05 on q/k/v/o; batch 16, context 1024; AdamW with learning rate 2×10^{-4} (adapters) and 10^{-3} (head), weight decay 10^{-4}, gradient clipping at 1.0.

Protocol. Holdouts are the last 10% of each dataset by position (∼3 h on 5 October; 14:30–17:33 for the sub-second set), purged and embargoed by the longest horizon the head forecasts. Confidence intervals use 15,000 evenly spaced, time-ordered holdout positions and a moving-block bootstrap [Künsch, 1989] with block length set to twice the longest horizon divided by the median spacing of the sampled positions, floored at 25 samples (the floor binds for the sub-second and 1/2/5 s sets; the 5/10/30 s set uses 108). Model-minus-baseline comparisons bootstrap the *difference* of the two correlations on identical resampled blocks: a paired bootstrap [Koehn, 2004], the analog for a non-decomposable statistic of a Diebold–Mariano comparison [Diebold and Mariano, 1995] in the spirit of White's reality check [White, 2000].

The 200 ms direction interval is the widest (±0.06) because the target is extremely heavy-tailed: excess kurtosis is ∼180, and the top 5% of $|r|$ windows carry 64% of the μ–r covariance, so the Pearson correlation inherits the sampling uncertainty of a few large-move blocks. The rank correlation is three times tighter (0.45 ± 0.02) and horizon-stable (0.43–0.45 at 0.2–1 s), and winsorizing the 1% tails *raises* the point estimate to 0.51 ± 0.02: the rarest, largest moves are

predicted less well than the bulk. We keep Pearson as the headline because it is the conventional, tail-inclusive statistic and the paired gaps are computed on it.

Sub-second results. Tables 4 and 5 and Figure 4 benchmark on an embargoed holdout. On *direction* the model beats order-book imbalance (OBI), including a 10-level book and OBI combined with recent-return momentum, at every sub-second horizon. The margin grows toward the short end and plateaus near +0.2 correlation from 0.5 s down to 200 ms. The 0.2–2 s gaps over OBI all exclude zero (Table 7 collects every paired gap); at 5 s the model and OBI are statistically tied (+0.03, CI $[-0.01, +0.08]$). A deeper book does *not* help: the 10-level imbalance is *worse* than the touch at 200 ms. Whatever the model reads is not available in a richer snapshot of the book.

Beyond 5 s the frozen representation ties OBI while the LoRA head degrades. At 10 and 30 s (Table 4, head C) the paired LoRA-minus-OBI gaps are −0.07 ($[-0.17, +0.04]$) and −0.22 ($[-0.44, +0.02]$), and the LoRA head's directional skill is gone by 30 s (-0.09 ± 0.19), whereas the frozen probe (0.15/0.12) still matches OBI (0.17/0.13). Beyond 5 s, then, the adapter stops adding to the frozen representation (on head C's window it is already below the probe at 5 s, 0.15 vs 0.19, Table 6); these long-horizon LoRA estimates are also the least precise in the table (±0.10, ±0.19), and we have not isolated the cause. It is not classical methods reclaiming direction: the representation itself keeps pace with OBI.

On *volatility* the same pattern holds but is thinner. The model beats GARCH and trivial persistence at 0.2–0.5 s by paired gaps of $+0.08/+0.09$ (GARCH) and $+0.06/+0.08$ (persistence), all CIs excluding zero, though the 0.2 s persistence gap only marginally ($[0.00, 0.16]$). At 1 s the sub-second head ties GARCH (−0.00, $[-0.04, +0.02]$) while still beating persistence (+0.07); the 1/2/5 s head already trails GARCH there (−0.12). From 2 s on GARCH wins significantly: −0.17 at 2 s and −0.09 at 5 s on head B's window, and $-0.30/-0.42/-0.12$ at 5/10/30 s on head C's, all CIs excluding zero. Persistence is the weaker opponent: on head B's window the model ties it at 1 and 2 s (−0.01, +0.02) and beats it at 5 s (+0.07, $[0.02, 0.12]$), and only on head C's window does persistence win, by $-0.14/-0.33/-0.17$ at 5/10/30 s, all CIs excluding zero. Beyond a second or two, volatility clustering dominates: its sufficient statistic, recent realized vol, is handed directly to GARCH but reaches the model only through 16-bin-quantized tokens.

These fine-tune correlations are not a seed artifact: re-run with a cosine-decayed schedule over four random seeds, the sub-second LoRA reproduces the values of Tables 4 and 5 within a tight seed spread (direction ±0.006, volatility ±0.03 at 0.2–1 s).

The target is the price, not the queue. At 200 ms the best bid or ask moves in only 27% of holdout windows (35% at 0.5 s, 45% at 1 s), so the plain mid $(P_b+P_a)/2$ is exactly unchanged 74%/66%/56% of the time and a microprice return in those windows is pure queue-size shifting. One might therefore ask whether the sub-second correlation measures price prediction or queue prediction. It is price prediction. Scoring the same μ_h against the *plain-mid* return gives correlation 0.52/0.50/0.42 at 0.2/0.5/1 s—*higher* than against the microprice it was trained on (0.44/0.45/0.39)—and 82%/83%/75% sign accuracy on the windows where the price actually moves (correlation 0.58/0.55/0.46 there). On the complementary windows where the touch does not move, μ_h has essentially no skill against the microprice change ($+0.07/-0.03/-0.09$): the model is not forecasting queue shifts. Because the microprice is itself built from touch imbalance, OBI is mechanically disadvantaged on that target; re-fitting OBI to the plain-mid return raises it to 0.35/0.31/0.28, yet the paired model-minus-OBI gap on the plain mid is $+0.17/+0.19/+0.14$ (95% CIs $[0.09, 0.24]$, $[0.16, 0.22]$, $[0.12, 0.17]$), indistinguishable from the paired microprice gap of Table 7 ($+0.22/+0.21/+0.16$). The edge survives on the plain price.

The fine-tuned model beats gradient boosting on order-flow features. OBI is a snapshot; the natural classical *flow* competitor is order-flow imbalance [Cont et al., 2014]—signed depth changes aggregated over a lookback—and its multi-level extension [Kolm et al., 2023]. We reconstruct the book and fit linear models on OFI at the top K=10 price levels over lookbacks matched to the horizon, $(t-h, t]$ and $(t-5h, t]$, trained on the train split and scored on the same embargoed holdout. The construction is sound—contemporaneous OFI–return correlation is +0.20—but its *predictive* power is weak: level-1 OFI alone gives 0.07/0.04/0.00 at 0.2/0.5/1 s, and ten-level OFI at the matched lookback reaches 0.23/0.20/0.16, the same level as the OBI snapshot (0.23/0.24/0.23); longer lookbacks are worse (0.14/0.11/0.09). A linear flow aggregate therefore captures what the snapshot captures and no more. A non-linear classical baseline, a gradient-boosted model on the same features (ten-level OFI at both lookbacks, OBI at one, five and ten levels, and the recent return), reaches 0.39/0.37/0.31, recovering most of the linear baselines' shortfall. The model's paired gap over it is $+0.05/+0.07/+0.08$ (95% CIs $[0.03, 0.08]$, $[0.05, 0.10]$, $[0.05, 0.10]$)—significant at every sub-second horizon, but roughly constant rather than growing toward the short end. The large, horizon-growing margins over OBI therefore partly measure how much a *linear* snapshot leaves on the table; against the best engineered-feature model, the fine-tuned model is worth a steady ∼+0.05–0.08 in correlation.

Table 4: Direction: correlation of prediction with the forward mid return on the embargoed holdout, as point ± half-width of a 95% moving-block-bootstrap CI (block bootstrap because overlapping forward windows autocorrelate the samples). "Scratch (17M)" is a same-tokens transformer trained from scratch, no pretraining (mean of 3 seeds). OBI (best) is the strongest of touch / 5-level / 10-level imbalance and their combination with recent-return momentum, each fit on train, the strongest chosen by holdout correlation (a choice that favors the baseline); trade-flow is signed trade volume. GBM is a gradient-boosted model on 10-level OFI (lookbacks h and $5h$), OBI at 1/5/10 levels and the recent return (HistGradientBoosting, 300 iterations, learning rate 0.05, 31 leaves, min leaf 200, fit on 400k train rows); computed for the sub-second set only. Bold marks the best point estimate only; comparisons use the paired gaps of Table 7. The *head* column says which LoRA produced the row: A = co-trained on 0.2/0.5/1 s, B = 1/2/5 s, C = 5/10/30 s, each on its own dataset and embargoed holdout; Table 6 shows the sensitivity to this grouping.

corr$(\cdot, r)$	head	OBI (best)	trade-flow	GBM	scratch (17M)	FM probe	FM LoRA
0.2s	A	0.23 ± 0.04	-0.12 ± 0.13	0.39 ± 0.04	0.19 ± 0.03	0.33 ± 0.04	$\mathbf{0.44} \pm 0.06$
0.5s	A	0.24 ± 0.04	-0.06 ± 0.07	0.37 ± 0.03	0.20 ± 0.03	0.38 ± 0.03	$\mathbf{0.45} \pm 0.03$
1s	A	0.23 ± 0.05	-0.03 ± 0.07	0.31 ± 0.03	0.20 ± 0.03	0.34 ± 0.03	$\mathbf{0.39} \pm 0.03$
2s	B	0.23 ± 0.03	-0.03 ± 0.04	—	0.17 ± 0.03	0.29 ± 0.03	$\mathbf{0.31} \pm 0.03$
5s	B	0.21 ± 0.05	0.04 ± 0.08	—	0.14 ± 0.04	0.23 ± 0.04	0.24 ± 0.03
10s	C	$\mathbf{0.17} \pm 0.05$	0.08 ± 0.09	—	0.09 ± 0.05	0.15 ± 0.05	0.10 ± 0.10
30s	C	$\mathbf{0.13} \pm 0.06$	0.09 ± 0.18	—	0.07 ± 0.07	0.12 ± 0.05	-0.09 ± 0.19

The edge is a short-horizon forecast. We probe how far ahead the signal reaches by sliding the realized window: we score μ_h against the return over $(t+\text{gap}, t+\text{gap}+h]$ as the lead gap grows (Fig. 5). The correlation is strongly front-loaded—it falls ∼60–80% within one horizon of lead and fades toward zero by two horizons, leaving only a small residual (clearest at the shortest horizon). The edge forecasts the *imminent* order-flow continuation rather than a persistent multi-horizon signal, which is why static snapshots and window-averaged scalar features miss it.

Table 5: Volatility: correlation of prediction with realized vol on the embargoed holdout (GARCH fit on train only, forecast out-of-sample), as point ± half-width of a 95% moving-block-bootstrap CI. "Scratch (17M)" is the from-scratch transformer (mean of 3 seeds; not scored for volatility on the 1/2/5 s set). *head*: A = LoRA co-trained on 0.2/0.5/1 s, B = on 1/2/5 s, C = on 5/10/30 s, each on its own dataset and holdout (Table 6). The FM interval widens at long horizons and is right-skewed at 10 s, where few independent windows fit the holdout (†: the 10 s FM entry; full CI $[0.19, 0.58]$); we do not read the non-monotone 5/10/30 s FM values as a horizon effect. Paired model-minus-baseline gaps (Table 7) are positive with CIs excluding zero at 0.2–0.5 s (vs GARCH $+0.08/+0.09$; vs persistence $+0.06/+0.08$) and a tie vs GARCH at 1 s on head A; on head B GARCH wins at 1–5 s while persistence ties at 1–2 s and loses at 5 s; on head C both baselines win at 5/10/30 s.

corr$(\cdot, rv)$	head	persistence	GARCH(1,1)	scratch (17M)	FM LoRA
0.2s	A	0.46 ± 0.09	0.43 ± 0.08	0.31 ± 0.06	$\mathbf{0.52} \pm 0.05$
0.5s	A	0.37 ± 0.07	0.37 ± 0.04	0.30 ± 0.06	$\mathbf{0.46} \pm 0.04$
1s	A	0.30 ± 0.05	$\mathbf{0.38} \pm 0.06$	0.29 ± 0.05	0.37 ± 0.06
1s	B	0.29 ± 0.05	$\mathbf{0.40} \pm 0.05$	—	0.28 ± 0.04
2s	B	0.26 ± 0.05	$\mathbf{0.45} \pm 0.05$	—	0.29 ± 0.06
5s	B	0.29 ± 0.07	$\mathbf{0.46} \pm 0.06$	—	0.37 ± 0.06
5s	C	0.49 ± 0.07	$\mathbf{0.65} \pm 0.06$	0.43 ± 0.07	0.35 ± 0.11
10s	C	0.61 ± 0.08	$\mathbf{0.70} \pm 0.05$	0.51 ± 0.07	$0.28 \pm 0.20^\dagger$
30s	C	$\mathbf{0.74} \pm 0.08$	0.69 ± 0.07	0.54 ± 0.07	0.57 ± 0.06

No leakage: fresh post-training data. The strongest check against a subtle leak is to score the trained model on data collected *after* training ends. On a dataset built from the latest order flow—beginning 130 min after the end of the fine-tuning dataset, holdout included—the model gives direction correlation 0.47/0.49/0.46 and volatility correlation 0.52/0.48/0.43 at 0.2/0.5/1 s: *at or above* the holdout. A leakage-inflated holdout would collapse here; instead it is reproduced. (The fresh window is quieter, so correlation rises while directional accuracy dips slightly—a regime effect, not leakage.) A second, structural check: the pretraining corpus of the main model was assembled on 2 October, before any of the fine-tuning data (3–5 October), so the fine-tune holdout lies strictly outside the pretraining window and no holdout sequence could have been memorized in pretraining (Fig. 6).

One adapter per horizon. Table 6 leaves an ambiguity: a head co-trained on 0.2/0.5/1 s and one co-trained on 1/2/5 s disagree about 1 s by ∼0.1, but they also differ in the window they are scored on. Either the co-training interferes, or the window does. We separate them by training *one adapter per target*—three LoRAs, each supervised on a single horizon—on the same backbone, the same fine-tuning set and the same split (the embargo is pinned to the joint runs' longest horizon, so the holdout is identical to the token), with the 20k-step cosine recipe and three seeds each, against a joint adapter trained identically. All are scored on one further window of 33.5 hours (15–16 October), disjoint from every fine-tuning set and later than every pretraining corpus used here.

Separating the horizons does not help; it costs a little. Direction falls from 0.465 to 0.453 at 0.2 s (seed ranges disjoint), and from 0.487 to 0.483 and 0.452 to 0.449 at 0.5 and 1 s (overlapping); volatility is a wash ($+0.006/-0.010/-0.006$, mixed signs). The loss is largest at the shortest and noisiest horizon, which is where auxiliary supervision should help most: the three targets are nested,

Table 6: Sensitivity to the horizon grouping. The same horizon scored by two different LoRA heads (each co-trained on its own horizon set, on its own dataset and ∼3 h embargoed holdout; the baselines move too because the holdout window differs). Direction at 1 s is stable; volatility at 1 s and direction at 5 s swing by ∼0.1—larger than the seed spread (±0.006 direction, ±0.03 volatility) and comparable to the claimed margins. Two confounds are bundled here, the co-trained horizons and the scoring window; the paragraph *One adapter per horizon* below rules out the co-training, leaving the window by elimination.

horizon	head	OBI	FM probe	FM LoRA (dir)	persistence	GARCH	FM LoRA (vol)
1 s	A (0.2/0.5/1)	0.23	0.34	0.39	0.30	0.38	0.37
1 s	B (1/2/5)	0.24	0.35	0.38	0.29	0.40	0.28
5 s	B (1/2/5)	0.21	0.23	0.24	0.29	0.46	0.37
5 s	C (5/10/30)	0.19	0.19	0.15	0.49	0.65	0.35

Table 7: Paired model-minus-baseline gaps in correlation (FM LoRA minus baseline) with 95% paired moving-block-bootstrap CIs on identical resampled blocks. The marginal intervals of Tables 4–5 overlap by construction, since every predictor is scored on the same windows and shares their block-to-block variation, so these paired gaps are the comparison statistics: an interval excluding zero means the ranking is reliable on this window, one spanning zero means a tie. *: interval excludes zero ($^\circ$: marginal, lower bound 0.000). Head as in Table 4; GBM was run for the sub-second set only.

horizon	head	dir: FM−OBI	dir: FM−GBM	vol: FM−GARCH	vol: FM−persistence
0.2s	A	$+0.22\ [0.17, 0.26]^*$	$+0.05\ [0.03, 0.08]^*$	$+0.08\ [0.02, 0.17]^*$	$+0.06\ [0.00, 0.16]^\circ$
0.5s	A	$+0.21\ [0.18, 0.24]^*$	$+0.07\ [0.05, 0.10]^*$	$+0.09\ [0.05, 0.12]^*$	$+0.08\ [0.02, 0.18]^*$
1s	A	$+0.16\ [0.14, 0.20]^*$	$+0.08\ [0.05, 0.10]^*$	$-0.00\ [-0.04, 0.02]$	$+0.07\ [0.01, 0.16]^*$
1s	B	$+0.14\ [0.12, 0.17]^*$	—	$-0.12\ [-0.16, -0.07]^*$	$-0.01\ [-0.06, 0.06]$
2s	B	$+0.08\ [0.05, 0.10]^*$	—	$-0.17\ [-0.23, -0.10]^*$	$+0.02\ [-0.02, 0.09]$
5s	B	$+0.03\ [-0.01, 0.08]$	—	$-0.09\ [-0.13, -0.06]^*$	$+0.07\ [0.02, 0.12]^*$
5s	C	$-0.03\ [-0.12, 0.06]$	—	$-0.30\ [-0.37, -0.17]^*$	$-0.14\ [-0.22, -0.01]^*$
10s	C	$-0.07\ [-0.17, 0.04]$	—	$-0.42\ [-0.52, -0.13]^*$	$-0.33\ [-0.43, -0.04]^*$
30s	C	$-0.22\ [-0.44, 0.02]$	—	$-0.12\ [-0.20, -0.06]^*$	$-0.17\ [-0.22, -0.11]^*$

so the longer ones act as a lower-variance teacher for the structure the adapter must expose—the standard multi-task argument, in the regime where it is expected to bite, since almost all of the adapter's capacity is shared and only a two-output layer is task-specific [Caruana, 1997, Baxter, 2000].

The useful consequence is negative: at most 0.012, an order of magnitude below the ∼0.1 swings of Table 6. Per-horizon training therefore *cannot* account for those swings, which leaves the scoring window as the explanation—consistent with the comparison across windows below, where absolute correlations move by 0.03–0.05 across windows while paired gaps stay put. We keep the joint head.

Across evaluation windows. A single ∼3-hour holdout cannot distinguish a real edge from a favorable afternoon, so the comparison is run on five holdouts: the 5 October window of Tables 4–7, adjacent to the fine-tuning data, and four drawn from disjoint, later stretches of the same instrument—8, 11, 14 and 17 October, of 7.1, 7.1, 7.5 and 12.5 hours, 34 hours between them.

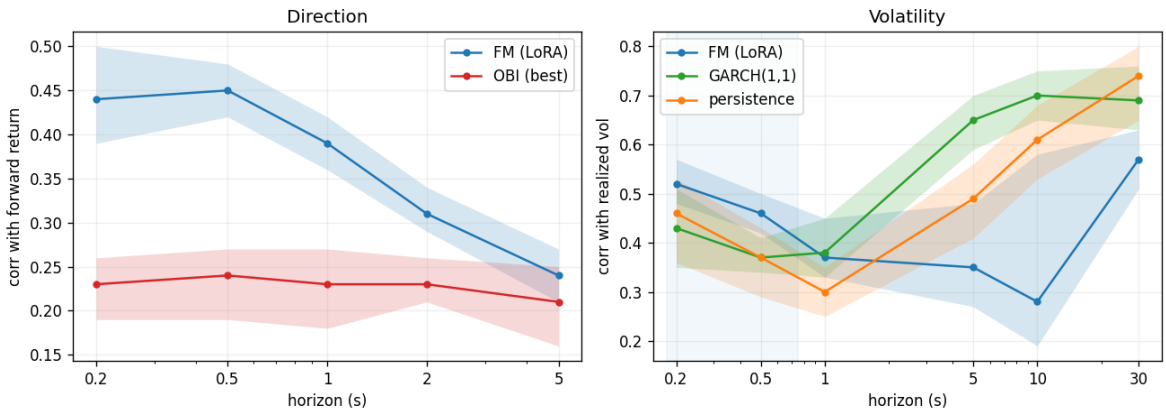

Figure 4: Sub-second results. *Direction* (left): the fine-tuned model beats the strongest order-book-imbalance baseline at every horizon, and the gap over this linear baseline widens toward the short end (over the gradient-boosted baseline of Table 4 the gap is a steady +0.05–0.08; not shown). *Volatility* (right): the model beats GARCH and persistence at 0.2–0.5 s (shaded) but is overtaken by $\geq$5 s, where the sufficient statistic (recent realized vol) dominates. Classical signals decay toward the ultra-short end, where the model's edge is largest. Direction points at 0.2–1 s are head A and at 2–5 s head B; volatility points at 0.2–1 s are head A and at 5–30 s head C, each head scored on its own holdout window (Table 6).

Each window is built, split and embargoed identically (last 10% by position, purged by the longest horizon, 15,000 evenly spaced positions, block bootstrap). Nothing about the model changes: the *same* adapter, fine-tuned once on 3–5 October, is carried forward to windows up to twelve days later, while OBI, the gradient-boosted model, GARCH and persistence are all *refit* on each window's own training split. The comparison is therefore stacked against the foundation model, which is stale by construction while every baseline is local and current. We do not refit the adapter per window: it costs a GPU fine-tune each time, and in preliminary experiments refitting on fresher data brought no improvement. We leave a full test to future work.

Table 8 reports the paired gaps. Direction holds without exception: the model beats order-book imbalance in 15 of 15 window-horizon cells and the gradient-boosted baseline in 15 of 15, every interval excluding zero. Against the non-linear baseline the margin on the four later windows is in fact *larger* than on 5 October—averaging +0.070/+0.113/+0.120 against +0.054/+0.074/+0.078—so the figure Tables 4–7 report is the conservative end of the range, not its center. Signed trade-flow remains uninformative everywhere (15 of 15 intervals spanning zero).

The volatility picture is weaker and, more usefully, *regime-conditional*: the model beats GARCH at 0.2 and 0.5 s on four of five windows but ties it on 17 October (+0.027 and +0.032, both spanning zero), and the tie at 1 s—reported as a null in Table 7—is itself the usual outcome (4 of 5). We therefore weaken the second-moment claim: the model improves on GARCH at the shortest horizons in most but not all market conditions.

Zero-shot cross-asset transfer. The backbone is pretrained on 73 asset streams across both venues (40 Bitstamp, 33 Bitfinex); ETH and SOL are among them on both venues, so the *representation* has seen their order flow, but the LoRA and head are fine-tuned on BTC alone. Applied *untouched* to Bitstamp ETH/USD and SOL/USD (correlation is scale-invariant, so the BTC-scale head is fair to compare), the BTC-trained model still beats each asset's own baselines: direction

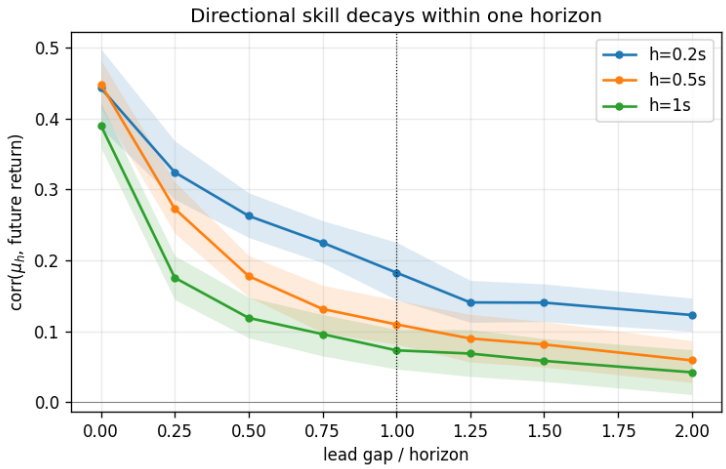

Figure 5: Directional skill vs lead time: corr(μ_h, return over $(t{+}\text{gap}, t{+}\text{gap}{+}h]$) against lead gap (as a fraction of h), with 95% moving-block-bootstrap bands. gap/h=0 is the trained target; the dotted line marks one full horizon of lead. The signal is a front-loaded, short-lived continuation.

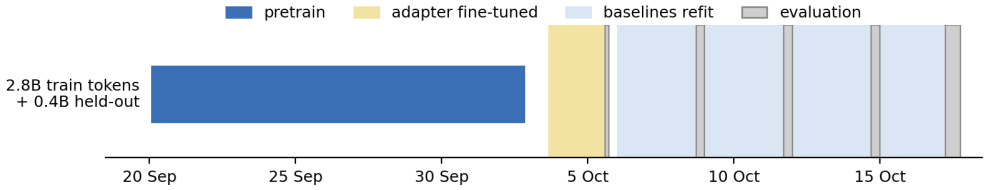

Figure 6: Fine-tuning and evaluation on the time axis. The bar is the backbone's pretraining corpus (training tokens plus the later held-out chunk), which ends before the fine-tuning data begins. The adapter is fine-tuned once, on 3–5 October (yellow), and the baselines are fit on the same span; the 5 October holdout that follows is the first evaluation window. Each later evaluation window of Table 8 is the last 10% of a three-day block by event count (so the quiet Saturday, 17 October, spans 12.5 h) whose first 90% (blue) refits the baselines alone; the adapter is carried forward unchanged. Model and baselines are scored together on the gray windows.

correlation 0.26–0.33 (ETH) and 0.13–0.32 (SOL) versus OBI ≤ 0.14 and ≤ 0.05 respectively, and volatility correlation ≥ 0.39 on both. Direction is $\sim$40% weaker than on BTC, and the 200 ms edge scales with liquidity (BTC > ETH > SOL), but at 0.5–1 s the thin SOL book transfers as well as ETH. The fine-tuned mapping is thus asset-general, with liquidity setting how short the horizon can be pushed—a strong form of the foundation-model claim, since the model predicts assets it was never fine-tuned for.

Does pretraining earn its keep? We ablate pretraining directly by training the *same architecture family end-to-end from scratch*: 17M parameters, random init, identical tokens and head, 3 seeds (Tables 4, 5). This gives a from-scratch model its best shot, not just an adapter on frozen random weights. The control trains on the *identical* supervised split (same targets, same embargo, $\sim$24 M labeled tokens) with *more* optimization than the LoRA fine-tune (up to 8,000 steps with early stopping vs. 1,500), so it is not short-changed on labels or compute. We size it at 17M because $\sim$24 M labeled tokens is already only $\sim$1.4 tokens per parameter; a 300M model trained end-to-end

Table 8: Paired gaps on five evaluation windows: the 5 October holdout of Tables 4–7 and four later, disjoint holdouts (7.1–12.5 h each; 34 h in total). The adapter is unchanged and fine-tuned once on 3–5 October; every baseline is refit on each window's own training split. Absolute correlations move across windows (the model's 0.2 s direction 0.390–0.509, OBI's 0.132–0.261) while the paired gap stays within 0.215–0.279. *: the 95% paired moving-block-bootstrap interval excludes zero.

gap	h	5 Oct	8 Oct	11 Oct	14 Oct	17 Oct
FM − OBI	0.2	+0.215*	+0.258*	+0.279*	+0.224*	+0.252*
	0.5	+0.213*	+0.240*	+0.250*	+0.202*	+0.231*
	1.0	+0.161*	+0.200*	+0.194*	+0.154*	+0.196*
FM − GBM	0.2	+0.054*	+0.049*	+0.089*	+0.039*	+0.103*
	0.5	+0.074*	+0.103*	+0.143*	+0.078*	+0.128*
	1.0	+0.078*	+0.122*	+0.140*	+0.097*	+0.119*
FM − GARCH	0.2	+0.084*	+0.119*	+0.115*	+0.075*	+0.027
	0.5	+0.085*	+0.072*	+0.083*	+0.082*	+0.032
	1.0	−0.004	−0.004	+0.044	+0.025	+0.059*
FM − persistence	0.2	+0.061*	+0.103*	+0.175*	+0.041	+0.183*
	0.5	+0.082*	+0.141*	+0.218*	+0.118*	+0.166*
	1.0	+0.074*	+0.102*	+0.167*	+0.089	+0.166*

on it would be far into memorization. A size sweep bears this out. From scratch, direction correlation *falls* with model size (0.20/0.13/0.05 at 1 s for 17/75/300M, all below OBI's 0.23), and at the FM's own 300M size it collapses to chance while pretraining+LoRA reaches 0.39: identical architecture and labels, pretraining the sole difference. What it lacks is the 2.76 B *unlabeled* pretraining tokens, the quantity under test.

On *direction* the from-scratch model learns only about *half* the fine-tuned model's correlation at every horizon (0.14–0.20 vs 0.24–0.45): a real signal, but most of it is a product of pretraining. It never beats even the simple order-book-imbalance baseline (Table 4: 0.14–0.20 vs OBI's 0.21–0.24) and sits ∼0.2 below the gradient-boosted baseline (0.39/0.37/0.31), whereas the pretrained model clears that baseline at every sub-second horizon. Pretraining is what moves the model from well below the best engineered-feature baseline to slightly above it.

On *volatility* the effect is horizon-dependent. Pretraining helps sub-second (0.31/0.30 scratch vs 0.52/0.46 at 0.2/0.5 s) but is irrelevant beyond 5 s, where the from-scratch model *matches or exceeds* it (0.43/0.51/0.54 vs 0.35/0.28/0.57), because long-horizon volatility is a sufficient statistic learnable by direct supervision. Pretraining is thus load-bearing for direction at every horizon and for volatility only below ∼1 s, concentrated exactly where the model wins. A cheaper control (probe/LoRA on frozen random weights) agrees: random-init direction is negligible (probe/LoRA corr 0.009/0.063 at 1 s on the 1/2/5 s set, against 0.38/0.40 for the pretrained probe/LoRA pair trained in the same ablation run; the 1,500-step runs of Table 6 give 0.35/0.38 for the same head and horizon).

Pretrain quality transfers to direction, not volatility. We fine-tune four backbones spanning a wide out-of-sample LM loss (NLL 2.94–3.43 scored as a language model on the 5 October fine-tune holdout; this is a third loss scale, since Table 2 scores each model on the later held-out chunk and Fig. 2 on a 12–13 October chunk, so only orderings carry across them), all with pretraining corpora assembled before any of the fine-tuning data, with an identical LoRA recipe (rank 16, cosine-decayed schedule, 10k steps). To separate a genuine backbone effect from training-seed vari-

ance we run four seeds, with the holdout held fixed, for the two endpoints: the weakest backbone and the paper's main model.

Seed variance is sharply moment-dependent. Directional correlation is reproducible to ± 0.006 across seeds, whereas volatility scatters by ± 0.03, so single-run comparisons of pretrain quality sit below the noise. With seeds, *direction* shows a clean, statistically unambiguous transfer: all four main-model seeds exceed all four weakest-backbone seeds at every sub-second horizon ($+0.027/+0.019/+0.016$ correlation at 0.2/0.5/1 s, four times the seed spread; one-sided permutation $p = 1/70 \approx 0.014$, the smallest value four seeds against four can give). Across the single-run curve (0.2 s) the gain is a step rather than a slope: the two weakest backbones sit at 0.40, the third at 0.42, the main model at 0.43. A fifth backbone (v5, NLL 2.77), excluded from the comparison because its corpus was assembled on 5 October and may contain the holdout, lands in the same place (0.43; seeded 0.431 vs 0.434), so a backbone that may have seen the holdout gains nothing from it: further evidence against leakage. *Volatility* transfer is not resolvable above the ± 0.03 seed scatter.

Better-pretrained backbones therefore give higher downstream directional correlation. We do not read a scaling law from four mixed-corpus checkpoints; we report only that *better pretraining indicates better downstream direction*. Why the gain is not larger has three candidate causes we cannot separate: a near-efficient-market ceiling, limited pretraining scale, and a single-asset, single-regime fine-tune set. Crypto order flow is not supply-limited the way language corpora are, so two are directly testable (a fuller same-corpus pretraining-scale study than the single control of Section 4, and a multi-asset fine-tune), which we leave to future work. A *frozen probe* does not capture even this: its downstream correlation is flat in backbone quality, so a richer backbone is exploitable through adapters but not through a linear read-out.

Predictable $\neq$ tradeable. None of the sub-second forecasts clear a $\sim$2 bps one-way cost (half the $\sim$1–2 bps spread plus a high-volume taker fee of order 1 bp; retail fee tiers are an order of magnitude higher, which only strengthens the conclusion)—the moves are sub-basis-point. The signal is a real, leakage-free, asset-general *prediction*, but not directional alpha; its plausible uses are in execution (anticipating imminent adverse selection) rather than taking. This barrier is one of move *magnitude*, not forecast accuracy: more or more-diverse data could raise the correlation but not the size of a sub-basis-point move, so the below-cost ceiling is unlikely to yield to scale even as the prediction improves.

Inference latency. An execution use is only meaningful if a forecast arrives well inside the horizon it describes, so we measure the whole path from a raw order-flow event to a prediction (Table 9; batch 1, bf16, context 1024, 200 timed repetitions after 30 warm-up, device-synchronized). Tokenization—the C++ hot loop that turns one event into one token, excluding file I/O, since a live feed delivers a parsed frame—costs 0.06–0.07 μs per event, six orders of magnitude below the forward pass, and the moment head costs 0.05 ms. The backbone is the whole budget. Two practical findings shape it. First, the adapter is free once *merged*: LoRA is an additive low-rank update, so folding it into the base weights is exact, and leaving it unmerged adds $\sim$144 small matrix products (four projections $\times$ 18 layers $\times$ two factors) worth 6 ms of pure kernel-launch overhead on the H100—36% of total latency for 0.7% of the weights. Second, batch-1 inference is launch-bound rather than compute-bound: the 1024-token forward is $\sim$0.6 TFLOP ($\sim$2 ms of H100 compute) yet takes 10.4 ms eagerly, and on the H100 a full re-encode and a single-token step against a KV cache cost the same (10.4 vs 10.2 ms), which is only possible if neither is doing arithmetic. CUDA-graph capture removes that overhead. We quote the *full re-encode* under capture, without a KV cache,

because the two could not be combined in our implementation; a streaming system does strictly less work per event, so the figure is an upper bound on its cost. That figure—4.0 ms, with the 99th percentile 0.1 ms above the median—is 2% of the shortest horizon we forecast and sustains 247 events/s against the 199 events/s median rate of this instrument, though not the ∼1,000 events/s bursts, where one would forecast on a subsample. These numbers bound the model's own cost on one GPU; they exclude a feed handler, an order gateway and the network, and so are not a tick-to-trade figure.

Table 9: End-to-end inference latency, event to forecast. "Merged" folds the LoRA into the base weights (exact); "graphs" is CUDA-graph capture via `torch.compile(mode="reduce-overhead")`. All rows re-encode the full 1024-token window except the last, marked "streaming", which advances a KV cache by one token; the two regimes cost the same on the H100 because both are launch-bound there. Tokenization (0.06–0.07 μs/event) and the head (0.05 ms) are included in the end-to-end column but are negligible against the backbone. Sustainable rate is the reciprocal of the median.

configuration	device	p50 (ms)	p99 (ms)	events/s
LoRA unmerged, eager	H100	16.4	17.2	61
merged, eager	H100	10.4	11.8	96
merged + graphs	H100	**4.0**	**4.1**	**247**
LoRA unmerged, eager	RTX 3090	25.5	30.3	39
merged, eager (streaming)	RTX 3090	12.4	18.2	81

7 Discussion

What the decoupling suggests. Across our checkpoints, a lower next-token loss buys a better event and book-factor generator but not a better open-loop generator of price paths: the frequency of deep-sweep rollouts shows no sign of falling as loss falls, while the same lower loss does transfer, through adapters, to a better fine-tuned forecast (§6). The representation improves with loss; the simulation does not. A plausible remedy on the generative side is structural—feeding the emergent mid-price change back into the context each step, a lever TradeFM uses and our tokens-only model omits—and we leave it to future work.

Limitations. Fine-tuning is on a single asset (bitstamp:btcusd) and a single two-day window, though we test transfer zero-shot on two other assets; and while the directional result rests on five disjoint holdouts totaling 37 hours (Table 8), all of it is one instrument on one venue within a four-week stretch, so regime coverage is narrow. The corpus is ∼2 weeks (one market regime), so the flat pretrain-to-probe transfer may be scale-limited rather than intrinsic. Sub-second targets are discrete and the absolute margins modest (∼+0.06 on vol); GARCH on 0.2 s intervals is outside its native habitat, and the comparison across windows shows the second-moment claim to be regime-conditional—the model ties GARCH at 0.2–0.5 s on one of the five windows. At long horizons the volatility comparison is regime-sensitive at the 3-hour-window scale: GARCH's own 5 s skill moves from 0.46 to 0.65 between two such windows (Table 6), so the size of its margin over the model should not be read as a stable quantity. On the generative side, rollout volatility is tail-driven and run-to-run noisy, and we do not claim L1 parity with TradeFM. These remain exploratory results, not a benchmark.

Future work.

1. Pretraining on more data. With a control for corpus recency, we established that more pretraining data leads to a more accurate fine-tuned model. At the moment we only pretrain on a few weeks' worth of data. Once we have collected more data, 100 days and more, we will retrain the models and test the resulting price and volatility prediction accuracy.

2. Per-asset fine-tunes. Here we fine-tuned the pretrained model to predict the first two moments for the BTC/USD pair. We tested the resulting model on other assets and found that it beats basic baselines there. To understand how much accuracy we are leaving on the table by not fine-tuning per asset, we should test our zero-shot transfer accuracy against a per-asset fine-tune.

3. Larger model. Our experiments indicate that a larger model has a modestly lower negative log-likelihood, hence scaling the model size might lead to a more accurate model and better downstream accuracy.

4. State of the order book. The model we train here only looks at the order flow within the 1024-token context window. The model does not know the state of the order book at the beginning of the window. The state of the order book is potentially useful information, hence encoding it and feeding it to the model might improve the accuracy.

5. A longer context window. For pairs with a lot of transactions, such as BTC/USD, a context window of 1024 tokens only covers a few seconds of the order flow. Since the model can only see 1024 tokens and nothing before that, we are truncating a lot of potentially useful information. Hence increasing the context window size might lead to higher accuracy.

6. Modifications to the architecture. Linear attention or sliding-window attention might be helpful to save on compute once we have a longer context window.

8 Conclusion

We have shown that a generative model pretrained to predict the next token in Level-3 order flow can be fine-tuned to predict future microprice and volatility. The fine-tuned model beats classical benchmarks in predicting the microprice change direction at sub-second horizons and transfers zero-shot from the asset it was fine-tuned on to others. Better pretraining transfers downstream for direction: as the pretrained model gets better, the downstream fine-tuned model gets more accurate as well, which suggests that a larger pretraining dataset may improve the price prediction model further. All this indicates that a next-token prediction model learns a representation that is useful for predicting price at sub-second horizons.

References

Fernando Acero. What do foundation models for market microstructure learn? A controlled study with synthetic limit order books. *arXiv preprint arXiv:2610.20104*, 2026.

Jakob Albers, Mihai Cucuringu, Sam Howison, and Alexander Y. Shestopaloff. The good, the bad, and latency: exploratory trading on Bybit and Binance. *Quantitative Finance*, 25(6):919–947, 2025.

Abdul Fatir Ansari et al. Chronos: Learning the language of time series. *Transactions on Machine Learning Research*, 2024.

Emmanuel Bacry, Iacopo Mastromatteo, and Jean-François Muzy. Hawkes processes in finance. *Market Microstructure and Liquidity*, 1(1), 2015.

Jonathan Baxter. A model of inductive bias learning. *Journal of Artificial Intelligence Research*, 12:149–198, 2000.

Tim Bollerslev. Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 31(3):307–327, 1986.

Rich Caruana. Multitask learning. *Machine Learning*, 28(1):41–75, 1997.

Rama Cont. Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative Finance*, 1(2):223–236, 2001.

Rama Cont, Arseniy Kukanov, and Sasha Stoikov. The price impact of order book events. *Journal of Financial Econometrics*, 12(1):47–88, 2014.

Fulvio Corsi. A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics*, 7(2):174–196, 2009.

Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. A decoder-only foundation model for time-series forecasting. In *ICML*, 2024.

Francis X. Diebold and Roberto S. Mariano. Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3):253–263, 1995.

Sean Foley, William Krekel, Vito Mollica, and Jiri Svec. Not so fast: Identifying and remediating slow and imprecise cryptocurrency exchange data. *Finance Research Letters*, 51:103401, 2023.

Edward J. Hu et al. LoRA: Low-rank adaptation of large language models. In *ICLR*, 2022.

Rakshit Jha, Mattijs De Paepe, Samuel Holt, James West, and Shaun Ng. Deep learning for digital asset limit order books. *arXiv preprint arXiv:2010.01241*, 2020.

Maxime Kawawa-Beaudan, Srijan Sood, Kassiani Papasotiriou, Daniel Borrajo, and Manuela Veloso. TradeFM: A generative foundation model for trade-flow and market microstructure. *arXiv preprint arXiv:2602.23784*, 2026.

Philipp Koehn. Statistical significance tests for machine translation evaluation. In *EMNLP*, pages 388–395, 2004.

Petter N. Kolm, Jeremy Turiel, and Nicholas Westray. Deep order flow imbalance: Extracting alpha at multiple horizons from the limit order book. *Mathematical Finance*, 33(4):1044–1081, 2023.

Hans R. Künsch. The jackknife and the bootstrap for general stationary observations. *The Annals of Statistics*, 17(3):1217–1241, 1989.

Junjie Li, Yang Liu, Weiqing Liu, Shikai Fang, Lewen Wang, Chang Xu, and Jiang Bian. MarS: A financial market simulation engine powered by generative foundation model. In *ICLR*, 2025.

Eljas Linna, Kestutis Baltakys, Alexandros Iosifidis, and Juho Kanniainen. LOBERT: Generative AI foundation model for limit order book messages. *arXiv preprint arXiv:2511.12563*, 2025.

Peer Nagy, Sascha Frey, Silvia Sapora, Kang Li, Anisoara Calinescu, Stefan Zohren, and Jakob Foerster. Generative AI for end-to-end limit order book modelling: A token-level autoregressive generative model of message flow using a deep state space network. In *ICAIF*, 2023.

Peer Nagy et al. LOB-Bench: Benchmarking generative models of limit order book data. *Preprint*, 2025.

Maximilian Seitzer, Arash Tavakoli, Dimitrije Antic, and Georg Martius. On the pitfalls of heteroscedastic uncertainty estimation with probabilistic neural networks. In *ICLR*, 2022.

Justin Sirignano and Rama Cont. Universal features of price formation in financial markets: perspectives from deep learning. *Quantitative Finance*, 19(9):1449–1459, 2019.

Eric Smith, J. Doyne Farmer, László Gillemot, and Supriya Krishnamurthy. Statistical theory of the continuous double auction. *Quantitative Finance*, 3(6):481–514, 2003.

Srijan Sood, Maxime Kawawa-Beaudan, Zhijin Guo, and Daniel Borrajo. One token per trade: Multi-resolution limit order book forecasting with a foundation model. In *ICML 2026 Workshop on AI Forecasting*, 2026.

Sasha Stoikov. The micro-price: a high-frequency estimator of future prices. *Quantitative Finance*, 18(12):1959–1966, 2018.

Hugo Touvron et al. LLaMA: Open and efficient foundation language models. *arXiv:2302.13971*, 2023.

Halbert White. A reality check for data snooping. *Econometrica*, 68(5):1097–1126, 2000.

Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Unified training of universal time series forecasting transformers. In *ICML*, 2024.

Zihao Zhang, Stefan Zohren, and Stephen Roberts. DeepLOB: Deep convolutional neural networks for limit order books. *IEEE Transactions on Signal Processing*, 67(11):3001–3012, 2019.